Public Opinion Quarterly (2026) **Vol 00 No 00,** 1–29

An Experimental Comparison of AI-Enabled Semi-Structured Interviews and Fixed Surveys: Respondent Satisfaction, Response Patterns, Quality, and Representation

Amanda Austin[1], Edward Hohe[2], Ryan Kennedy[3,*], Leib Litman[4], William Minozzi[5], Laura Moses[6]

[1]PhD Candidate, Department of Political Science, University of Houston, Houston, TX, US
[2]PhD Candidate, Department of Political Science, The Ohio State University, Columbus, OH, US
[3]Professor, Department of Political Science, The Ohio State University, Columbus, OH, US
[4]Professor, Department of Psychology, Lander College, Kew Gardens Hills, NY, US, and Chief Research Officer, Cloud Research, New York, NY, US
[5]Professor, Department of Political Science, The Ohio State University, Columbus, OH, US
[6]Research Director, Institute for Democratic Engagement and Accountability (IDEA), The Ohio State University, Columbus, OH, US

Abstract Artificial intelligence (AI) interviewing offers the potential to enhance the depth and quality of survey research, yet it may also introduce respondent fatigue and unintended consequences for the measurement of political attitudes. From the standpoint of respondents, more extensive interviews impose greater cognitive and temporal burdens, which may reduce engagement or increase satisficing. From the perspective of political behavior researchers, deeper probing may shift the subset of beliefs or arguments that respondents access when forming opinions, thereby altering the distribution and apparent stability of measured attitudes. To investigate these dynamics, we conduct a large-scale, fully randomized experiment. We find that AI-mediated interviewing significantly increases both the volume and specificity of respondents' articulated rationales, but any decline in participant satisfaction is minuscule and we detect no increase in participant rolloff. However, AI interviews appear to shape subsequent responses: After offering justifications through AI interaction, respondents exhibit greater polarization in follow-up items. This result suggests that AI interviewing may alter the cognitive architecture of opinion formation. These findings both highlight the potential for

*Corresponding author: Ryan Kennedy, Department of Political Science, The Ohio State University, Columbus, OH, US; email: kennedy.310@osu.edu.

© The Author(s) 2026. Published by Oxford University Press on behalf of American Association for Public Opinion Research.
This is an Open Access article distributed under the terms of the Creative Commons Attribution License (https://creativecommons.org/licenses/by/4.0/), which permits unrestricted reuse, distribution, and reproduction in any medium, provided the original work is properly cited.
https://doi.org/10.1093/poq/nfag063

Downloaded from academic.oup.com/poq/advance-article/doi/10.1093/poq/nfag063/8762750 by guest on 10 November 2026

AI interviewing and underscore the need for studies of how AI interviewing may become a source of attitudinal change.

> "Since the beginning of social research, students have tried to combine the value of detailed qualitative applications with the advantages of more formalized techniques which could be managed on a mass basis" (Lazarsfeld 1944, 59).

The incorporation of AI in public opinion research is widely regarded as transformational, particularly for enabling in-depth interviews at scale. By providing a structured exploration of the cognitive processes underlying attitudes, AI offers researchers the means to better understand the foundations of political opinions. AI interviewers thus have the potential to advance our understanding of opinion formation and reasoning in mass publics (e.g., Xiao et al. 2020; Chopra and Haaland 2023; Wuttke et al. 2024). Already, various AI interviewing approaches are being evaluated (Barari et al. 2025), and AI interviewers are being integrated into experimental designs (Velez and Liu 2025) and used to develop adaptive Likert-style questions (Velez and Liu 2025).

While AI interviewing promises to enhance the depth and quality of survey responses, it may also have unintended consequences. Practically, more in-depth interviews demand greater time and cognitive effort from participants, resulting in substantially longer, more costly responses (Geiecke and Jaravel 2024). Such demands may impact satisfaction and completion rates, potentially introducing new biases in selection and attrition.

AI interviewing, moreover, may destabilize both attitude measurement and response interpretation. In-depth interviewing often prompts participants to provide reasons, perhaps even transforming their opinions in the process, not unlike public deliberation (Thompson 2008; Kalla and Broockman 2020; Minozzi et al. 2024). As such, this new interview mode itself may constitute a stimulus, altering the accessibility and salience of considerations. AI interviewing thus puts distance between traditional top-of-the-head responses to closed-ended questions and summary responses preceded by AI interviews, leaving researchers with no clear answer about which constitutes a "true" attitude.

We address these questions through a large-scale survey experiment ($N = 2{,}243$), manipulating the presence and timing of an AI interviewer. We evaluate effects on reasoning, satisfaction, duration, item nonresponses (rolloff), and response consistency.

We document that AI interviewing substantially increases both the quantity and depth of respondents' reason-giving, and that these effects come at the cost of additional completion time, as expected. Despite

Downloaded from academic.oup.com/poq/advance-article/doi/10.1093/poq/nfag063/8762750 by guest on 10 November 2026

these additional costs, AI interviewing has negligible effects on participant engagement and satisfaction, and we detect no effect on rolloff. AI interviewing also alters responses. Responses to follow-up (closed-ended) summary questions are more polarized for participants who responded first to a closed-ended question and then explained their reasoning in an AI interview. AI interviewing itself may, therefore, actively shape cognitive processing in ways that need further exploration.

AI Interviewing and Survey Respondents

The emergence of AI-driven survey tools prompts the question of how and why AI interviewing impacts respondent behavior. Technological innovations in survey administration have consistently presented both opportunities and concerns for researchers. For example, the move from face-to-face to digital interfaces had a range of impacts, both positive and negative. Digital interfaces reduced social desirability bias, especially around sensitive topics (Tourangeau and Yan 2007; Kreuter et al. 2008), but this came at a cost of lower response rates and potential nonresponse bias (Manfreda et al. 2008; Daikeler et al. 2020). The inability to clarify questions and differences between visual and auditory processing was also found to increase problems in interpretation and respondent satisficing (Heerwegh and Loosveldt 2008).

AI interviewing presents research questions that, in some ways, mirror these earlier technology questions and raise new ones. We develop several preregistered research questions (RQs) to explore AI interviewing's impact.[1]

From a respondent's perspective, participants are likely to spend more time interacting with an AI than they would with a standard set of closed-ended or fixed open-ended questions. As a result, respondents to AI interviews may find the experience less satisfying compared to more straightforward, less time-consuming and mentally taxing survey formats, especially if they are paid survey takers on a crowdsourcing platform. Research on open-ended questions has suggested that the increased cognitive burden results in greater rolloff and less satisfaction with surveys (e.g., Reja et al. 2003; Galesic 2006). Given that AI interviewing involves repeated open-ended questions with follow-ups, this literature would suggest substantial problems in respondent satisfaction.

On the other hand, survey respondents may value the opportunity to express their opinions in greater detail, no longer constrained by predefined response options. AI has proven effective in sparking meaningful

1. We preregistered all research questions and treatments at https://osf.io/4nsb8/overview prior to fielding our study. See Supplementary Material section 6 for discussion.

Downloaded from academic.oup.com/poq/advance-article/doi/10.1093/poq/nfag063/8762750 by guest on 10 November 2026

conversations (Croes and Antheunis 2020; Yin et al. 2024), and studies suggest that people enjoy explaining *why* they behave or think a certain way, driven by a desire to be heard and understood (Hardin and Higgins 1996; Brooks and John 2018; Hart et al. 2021). These potentially opposing possibilities led us to our first research question:

RQ1. Will respondents be more or less satisfied with the survey experience if it includes an AI interaction?

A second key question is whether the inclusion of AI interviews significantly influences the opinions respondents express. The most widely accepted question-response model in public opinion research identifies four stages in a response: comprehension, retrieval, judgment, and response (Tourangeau et al. 2000; Willis 2004). Comprehension refers to how the respondent interprets the question; retrieval is the process of recalling relevant information from memory; judgment involves evaluating that information and deciding how to answer; and response is the act of selecting an answer based on the available options, influenced by social desirability, question format, and the survey context. Similarly, Zaller (1992) introduces the concept of "sampling" in opinion expression, which aligns with retrieval and judgment, suggesting that individuals base their responses on the most recent and accessible considerations in their memory.

The potential impact of AI interviews on responses may be most apparent in terms of an ordering effect. There are several reasons to believe that the inclusion and timing of AI interviews could influence answers to closed-ended questions. For one, interacting with an AI interviewer may alter the cognitive processes involved in answering. If the interview begins with an AI interviewer, it may trigger fundamentally different thought patterns than those involved in simple retrieval. The process of thoughtfully articulating and justifying one's reasoning may influence the way respondents ultimately answer. Some scholars have suggested that this "deliberation within" can significantly shape the opinions individuals express (Goodin 2000; Goodin and Niemeyer 2003). Alternatively, however, research on "think aloud" protocols in survey research suggests that verbalization is more about justifying decisions made before cognitive processes (Lodge and Taber 2013). As such, leading with an AI interviewer may or may not produce different responses than immediately starting with a closed-ended question.

RQ2. Will responses to closed-ended questions vary if they are preceded or followed by an AI interview?

The cause-and-effect relationship suggested by RQ2 may also be reversed: Starting with a closed-ended summary question could trigger different behavior in a follow-up AI interview compared to beginning with the AI

Downloaded from academic.oup.com/poq/advance-article/doi/10.1093/poq/nfag063/8762750 by guest on 10 November 2026

interview itself. Respondents who first answer closed-ended questions and then engage in an AI interview might approach the interaction differently from those who start with an AI interview. People often rationalize their previously expressed opinions, retrieving information mainly to confirm their prior responses (Vedejová and Čavojová 2022). While this heuristic generally applies to information search and acceptance, a similar process is possible in a communication context. In this case, the respondent may be more inclined to justify their closed-ended answer, rather than engage in a deeper exploration of their thoughts on the issue.

RQ3. Do responses to an AI interviewer depend on whether it is preceded or followed by a closed-ended summary question?

Interaction with an AI interviewer may also serve to solidify—or even polarize—attitudes, as individual survey responses are often unstable, even within the same survey, due to variations in the sampling process (Feldman 1989). Engaging with an AI interviewer prior to answering a closed-ended question might enhance the consistency of responses by reinforcing the information retrieved earlier, making it more likely that the sampling process will remain stable over time. On the other hand, if the AI interview follows an initial closed-ended question, it could prompt some level of internal deliberation, potentially leading to shifts in opinion. These changes could manifest when the closed-ended question is asked again later in the survey. The direction of this shift, however, is uncertain. Deliberation could moderate attitudes by bringing in counter considerations more frequently (Tsai et al. 2024), but it may also provide more room for motivated reasoning and confirmation bias, potentially leading to more polarized responses (Taber and Lodge 2006). This possibility leads to two additional research questions:

RQ4. Will interaction with the AI interviewer alter response consistency?
RQ5. Will interaction with the AI interviewer polarize responses?

Following up closed-ended questions with open-ended probes is not a new practice. In fact, much of the potential of AI interviews appears to build on what can already be achieved with fixed, nonresponsive, open-ended follow-up questions. Therefore, an important question is whether, and to what extent, AI interviewers can provide responses that are different from those offered by traditional methods. Beyond simple measures like word count, AI interviewers have the ability to probe more deeply into respondents' reasoning, as their follow-up questions can be tailored to the specific answer provided. This customization may allow AI to gather more detailed insights into the reasons behind a respondent's answers. Additionally, AI may be capable of identifying and addressing any misinterpretations of questions, offering respondents a chance to

Downloaded from academic.oup.com/poq/advance-article/doi/10.1093/poq/nfag063/8762750 by guest on 10 November 2026

refine their answers. If AI interviews can yield deeper insight into respondents' thought processes—both in content and reasoning—they could substantially enrich the quality of survey data.

RQ6. How do responses to AI interviewers differ, if at all, from those to standard, fixed, follow-up open-ended questions?

Finally, we also preregistered two additional questions regarding heterogeneity of responses based on respondent attention, as well as demographic and personality characteristics. These results are omitted from the main text for space reasons, but are available to the interested reader in Supplementary Material sections 10 and 11. No issues of heterogeneity were noted that would impact the interpretation of the findings herein.

Data and Methods

To answer these research questions, we used a fully randomized design to evaluate the effects of AI interviewing. We elected to conduct this research within the Engage platform, designed by CloudResearch.[2] We opted to use this platform because: (1) It is one of the better-tested AI platforms, having been used throughout the 2024 election cycle in collaboration with the Siena Poll (e.g., Siena College Research Institute 2024); and (2) it features a robust design and deployment environment with multiple guardrails to ensure interview quality and data security, while avoiding issues like AI "hallucinations." Rather than relying on a single LLM, Engage employs different LLMs for specific modular components. For example, one LLM may design an initial response question, while another will check to make sure this follow-up question is in line with the interview prompt before posting it to the user. We used version 0.2.5 of Engage, which was deployed on October 10, 2024. While there is an infinite variety of behaviors, styles, and prompts that can be used with AI interviewing (e.g., Tsai et al. 2024), using a well-developed and easily accessible platform that has already incorporated a multi-agent structure for quality control ensures that our choices will not be too out of line with the expected mainstream of future research.

Engage is currently free for researchers, but, even if it were not, the total cost for conducting these interviews based on the number of tokens involved, and assuming a very inefficient system of implementation, would have been around $0.008 per interview, for a total of about $10.44

2. This experiment was reviewed by the IRB at The Ohio State University: Study #2024E1130, November 4, 2024.

Downloaded from academic.oup.com/poq/advance-article/doi/10.1093/poq/nfag063/8762750 by guest on 10 November 2026

for all the interviews.[3] Put simply, the implementation costs of AI interviews were dwarfed by participant costs.

At the time of this study, Engage did not have the ability to randomize treatment assignment. To overcome this issue, we designed each treatment arm separately in Engage and embedded the Engage survey links into the HTML of a Qualtrics survey.[4]

Independent Variables

Our experiment had five treatment arms, based on the components outlined in table 1. Table 2 records how the components were combined to create the five treatment conditions.

Four of these treatment arms included two closed-ended items authored by Pew Research Center (2024, 2018) that gauged support for recent immigration and tariff policy. The control condition, **Closed Only (CO)**, asked only the closed-ended questions without further discussion. A second non-AI treatment, **Closed, then Fixed (CF)**, included the same closed-ended items followed by a generic open-ended prompt: "Why do you hold this position on this issue?" Although more sophisticated techniques exist for fixed follow-up questions, such generic prompts are common in applied research.

The other three treatment arms all incorporated AI at varying stages. The **Closed, then AI (CAI)** follow-up condition entailed an open-ended AI interview after each closed-ended question. The **AI, then Closed (AIC)** follow-up arm reversed this order, starting with a general preamble asking what they think about the issue (O1 and O2), followed by the AI interview, and closing with the closed-ended questions. Finally, we included an **AI Only (AIO)** condition, which featured only the introductory preambles, followed directly by the AI-conducted interviews. Figure 1 shows an example of what a treatment (CAI) looked like in a simulated interview on the Engage platform.

Dependent Variables

We measured multiple outcomes to systematically assess our research questions. First, we assessed respondent behavior in the survey itself, including the time respondents took to complete the survey (in seconds) and the total number of words used in their text interview responses.

3. Calculations based on pricing for OpenAI's GPT-4o, the flagship model at the time of the experiment, on March 7, 2025. Details in Supplementary Material section 5.

4. The full survey text and HTML for Qualtrics integration are available in Supplementary Material section 2.

Downloaded from academic.oup.com/poq/advance-article/doi/10.1093/poq/nfag063/8762750 by guest on 10 November 2026

Table 1. Components used to construct treatments. Table shows the components that are used in the construction of treatments in table 2. AI portion shows the prompt used to generate the conversations, as illustrated in figure 1.

Component	Question/prompt
Closed-ended questions	
C1	How much do you agree or disagree with the following statement? There should be a way for undocumented immigrants to stay in the US legally if certain requirements are met. [Strongly agree, Agree, Somewhat agree, Neither agree nor disagree, Somewhat disagree, Disagree, Strongly disagree]
C2	Overall, do you think increasing tariffs or fees on goods imported from trading partners will be good or bad for the United States? [Very good, Good, Somewhat good, Neither good nor bad, Somewhat bad, Bad, Very bad]
Fixed open-ended question	
F1	Why do you hold this position on this issue?
Prompt for AI interview	
AI	Ask follow-up questions to develop a fuller understanding of why the respondent does or does not think [issue], or which way they lean if they say they neither agree nor disagree. Probe for specific reasons they hold their opinions. Assess the strength of their opinion and if they are open to counter-arguments that might change their thinking on the policy. Listen for and probe on: Language or phrasing, Emotion, Tone, their thoughts and feelings regarding policies towards [issue]. Continue with the interview until you get a FULL picture of what this person was thinking. Ask for clarification or more details when responses are vague or general. Ask a maximum of three follow-up questions! Don't ask leading questions, let participants spontaneously tell you what they think without putting words in their mouths. Also, always ask only one question at a time!
Open-ended opener for AI interview	
O1	What do you think of the idea that undocumented immigrants should be allowed to stay in the US legally if certain requirements are met?
O2	What do you think about increasing tariffs or fees on goods imported from trading partners to the United States?
Preamble to closed-ended question (following interview)	
P1	Now that you have discussed this and thought it through, please give us a summary of your thoughts on the issue.

Downloaded from academic.oup.com/poq/advance-article/doi/10.1093/poq/nfag063/8762750 by guest on 10 November 2026

Table 2. Composition of treatments. Table shows the composition of treatments using keys from table 1.

Treatment	Composition
Closed only (CO)	C1 and C2
Closed then fixed follow-up (CF)	(C1 + F1) and (C2 + F1)
Closed then AI follow-up (CAI)	(C1 + AI) and (C2 + AI)
AI then closed follow-up (AIC)	(O1 + AI + P1 + C1) and (O2 + AI + P1 + C2)
AI only (AIO)	(O1 + AI) and (O2 + AI)

For four experimental arms (excluding AIO), we analyzed responses to closed-ended items, both when they were asked in the Engage interview (which we label the "Interview" responses) and a second time after the demographic and personality questions ("Post-Interview" responses), allowing us to assess the consistency of respondents' answers. We also measured polarization by calculating the distance of each response from the median option and then taking the difference in these distances, such that positive values indicate polarization, negative values indicate moderation, and zero indicates no change in response. Additionally, we asked all respondents to rate their satisfaction with the survey experience on a seven-point scale and to describe their experience with three open-ended questions asking: (1) what they liked about the survey, (2) what they disliked about the survey, and (3) if they had any additional comments for the researchers. We coded these open-ended questions manually to record sentiment toward the AI interaction, as well as what other parts of the survey the respondents liked or disliked. Finally, we also used AI tools to measure the content of the text responses in Engage, specifically in terms of the rationales given by respondents for their opinions on immigration and tariff policies.

Covariates

We measured a number of respondent characteristics, to assess balance across experimental conditions. We drew standard demographics from CloudResearch: age, education, gender, partisanship, racial/ethnic identity, political interest, and political knowledge.

Additionally, we included two standard psychometric batteries: the Big-5 personality index (short form), which measures extroversion, agreeableness, conscientiousness, neuroticism, and openness (Herzberg and Brähler 2006), and a short-form measure of respondents' need for cognition, which may relate to their willingness to engage in thoughtful conversation on an issue (Cacioppo and Petty 1982). To assess covariate balance,

Downloaded from academic.oup.com/poq/advance-article/doi/10.1093/poq/nfag063/8762750 by guest on 10 November 2026

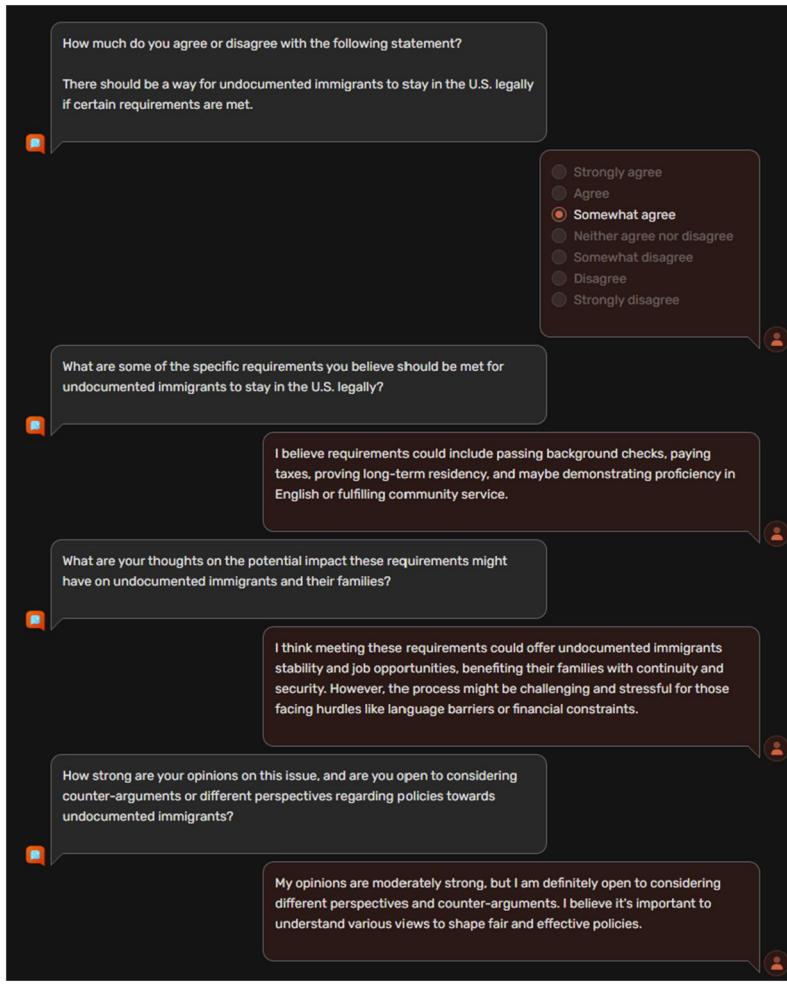


Figure 1. Example (simulated) of CAI treatment arm in the Engage platform. Figure shows a simulated conversation between AI agents on the immigration question.

we used equivalence testing (Hartman and Hidalgo 2018), which employs two one-sided tests to reject a null of meaningful difference and provide statistical evidence for equivalence between treatment groups. We found no imbalance across any treatment pair for any covariate.[5]

5. All equivalence tests were based on the "strict" null of a difference of 0.37 SD, using a cutoff of $p = 0.05$ (two-tailed), and were corrected for false discovery rates.

Downloaded from academic.oup.com/poq/advance-article/doi/10.1093/poq/nfag063/8762750 by guest on 10 November 2026

Recruitment

We recruited 2,243 respondents from CloudResearch's Connect platform on December 6–11, 2024, who were randomly assigned to one of the five treatment arms—approximately 450 participants per treatment (see Appendix A for details). CloudResearch has been used in a large number of experimental studies, replicates experimental findings from other major survey sources, and provides better data than many alternative online providers (Douglas et al. 2023; Hauser et al. 2023). The resulting sample was a nonprobability convenience sample on an opt-in basis (see Hartman et al. 2023). As an opt-in panel, Connect samples may differ from the US population on dimensions correlated with panel membership or survey-taking propensity. However, since our primary goal is causal inference from randomized treatment, internal validity should be maintained despite nonprobability recruitment. Furthermore, because participants from CloudResearch are used so often within academic research, they are typical of online participants in general.

To improve external validity, we took advantage of the option for CloudResearch to match the age, gender, ethnicity, and race characteristics of our sample against US Census values. However, to complete the survey in a timely fashion, we had to relax these quotas toward the end of data collection. This adjustment produced slight shortfalls in male, Hispanic/Latino, and white respondents, along with a larger shortfall in respondents over 60 years old. These quota fulfillment ratios are reported in Supplementary Material section 3.

Each respondent was paid a flat rate of $3 for their time, regardless of the treatment arm to which they were assigned. We estimated that the survey would take about 10–15 minutes for the most intensive treatment arm based on early pilots. The actual median response times ranged from about 6 minutes for the CO survey to about 14 minutes for the AIC survey.

Analysis

Unless otherwise noted, we utilized basic Ordinary Least Squares (OLS) regression for statistical tests. Results are generally presented as marginal means, with 67 percent and 95 percent confidence intervals. These values were estimated through 4,000 simulated draws from the variance-covariance matrix (King et al. 2000). Results in this main paper are presented for participants who passed the attention check unless otherwise noted. This decision, however, makes no difference in the conclusions.[6]

6. See Supplementary Material section 10 for more discussion.

Downloaded from academic.oup.com/poq/advance-article/doi/10.1093/poq/nfag063/8762750 by guest on 10 November 2026

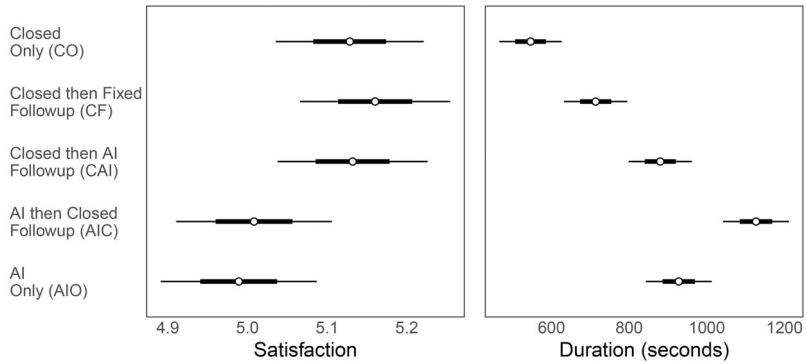


Figure 2. Negligible loss in satisfaction, despite increased duration. Satisfaction is measured on a seven-point scale; duration, in seconds. Points show the estimated marginal means. Bars show the quantile intervals for the 67 percent and 95 percent credible intervals estimated from 4,000 draws from the posterior distribution.

Survey Cost, Satisfaction, and Duration

A primary question for this research was how respondents would react to the greater cognitive burden and length involved in AI interviewing. RQ1 asked whether respondents would respond negatively to a task requiring more time and effort to receive the same level of compensation, or if they appreciated the opportunity to further express themselves.

To answer this question, we analyzed self-reported satisfaction with experiences on our survey. Figure 2 (left panel) displays marginal means for each treatment arm (i.e., the expected value for each treatment group). Satisfaction was generally higher in the control Closed Only (CO) and Closed then Fixed (CF) conditions, which did not include AI interviewing, but was also similar in Closed then AI follow-up (CAI), which did incorporate AI. Satisfaction was lower in both AI Only (AIO) and AI, then Closed follow-up (AIC) conditions, and some cross-treatment comparisons yielded significant differences. The scale of these differences, however, was very small—a little over 0.1 points on a seven-point scale. To formally test the difference due to AI specifically, we estimated bivariate linear regression models of satisfaction on an indicator of whether treatment included an AI component. On this seven-point scale, the coefficient was negative, at -0.06, with 95 percent interval $[-0.11, 0.02]$ for the full

Downloaded from academic.oup.com/poq/advance-article/doi/10.1093/poq/nfag063/8762750 by guest on 10 November 2026

sample.[7] We conclude that there is—at most—evidence of a very small decline in satisfaction due to AI interviewing.

The absence of satisfaction costs did not seem to be an artifact of differential attrition (i.e., rolloff). Of the 2,253 participants we recruited to start the survey, only 10 failed to complete it, and all of these respondents rolled off before starting the Engage component of the survey, which included all AI interactions. There was, however, a small level of differential completion of the text-entry components. Further inspection of this pattern revealed that it was almost certainly a technical issue dealing with when the "next" button is displayed.[8]

The negligible difference in satisfaction appears even though there were sizable, uncompensated differences in survey duration (see figure 2, right panel). As noted above, all respondents received $3 regardless of treatment. Duration varied dramatically across treatments, ranging from a median of less than 7 minutes for the Closed Only control (CO) to 14.5 minutes for an AI, then Closed follow-up interaction (AIC). Translating this into an hourly rate for interpretation, the result would be a median hourly rate of $25.90 for CO, $20.49 for CF, $14.59 for CAI, $13.90 for AIO, and $12.41 for AIC.

These differences in duration stand in stark contrast to the negligible differences in satisfaction. One explanation is that these respondents were professionals who were keenly aware of pay scales and thus unsurprised and undisturbed by any exigencies of participation. But another explanation is that the experience with the AI interviewer was sufficiently novel and interesting as to be intrinsically motivating.

Some evidence for the latter explanation emerged from qualitative data. Respondents were asked at the end of the survey what they liked best, what they disliked most, and if they had any additional comments for the researchers. We analyzed these comments manually. Despite no requirements for response, 1,057 of 1,422 respondents who interacted with an AI interviewer left such a comment (74 percent). Of those, 470 (44 percent) were positive. These included comments like "I enjoyed interacting with the chatbot because it made me think about my perspective," "The chatbot was interactive and interesting," "The interview with AI let me be myself and forget I'm taking a survey," and "I liked that chatbot section, mostly because it just gives me the chance to explain my thoughts rather than select some bubbles." The most common positive feedback described the conversation with the AI interviewer with words like "fun," "interactive," and "engaging."

7. Results for this comparison and all subsequent ones were similar when we included those who fail the attention check.

8. Supplementary Material section 8.2 provides a full breakdown of this issue.

Downloaded from academic.oup.com/poq/advance-article/doi/10.1093/poq/nfag063/8762750 by guest on 10 November 2026

Only 126 (12 percent) left some type of criticism about the AI. These included criticisms like, "Bots are repetitive and ask too many "yes, and' follow-ups," "I don't like AI in general because of the ethics surrounding how they are trained," and "The AI was a bit awkward and stilted in its language. It didn't feel like a real conversation." Even among the 126 respondents who left a negative comment, however, 47 (37 percent) leavened their criticism with a positive comment about the chatbot, for example, "I wish the chat with the chatbot could have gone on a little longer" or "It would have been nice to have a longer discussion with the chatbot." Overall, for every negative comment about the AI interviewer, there were six positive comments.[9]

While novelty might depress survey enjoyment by increasing cognitive demands, it might also elevate the survey experience by breaking up the monotony of what one respondent labeled "bubble hell surveys." Another respondent suggested they were able to offer more authentic, higher-quality data: "I liked answering questions through a conversation. Your [*sic*] getting 100 percent real honest whole answers that way, instead of having little bubbles to click"; and another said they liked the AI interview because it "made me think outside the box." These qualitative data support the latter interpretation. Those who mentioned that the survey is new, unique, or novel did so exclusively positively. And one respondent noted that this was their favorite AI survey, suggesting the satisfaction benefits of novelty might endure at least beyond the first couple of experiences.

We conclude that there is little indication that incorporating an AI interview into the survey results in lower survey satisfaction, even when it increases workload without additional payment.[10] Indeed, this evidence is consistent with the possibility that many respondents experienced intrinsic *enjoyment* because of the inclusion of the AI interviewer.

Word Count and Reason Provision

This section explores what unique insights interviews can reveal that cannot be uncovered through a simple open-ended follow-up—such as asking respondents why they answered as they did—and how the integration of AI into the interview process influences these outcomes.

9. A more complete breakdown of the negative comments is available in Supplementary Material section 8.3. All respondent comments are available in replication data.

10. Even the lowest pay amounts in our survey, however, were still within the bounds of what is usually considered "good" for an online survey. CloudResearch Connect provides a compensation rating system (https://connect.cloudresearch.com/); $12 an hour is considered "good" compensation in this rating. Our median compensation was above this threshold for all treatment arms.

Downloaded from academic.oup.com/poq/advance-article/doi/10.1093/poq/nfag063/8762750 by guest on 10 November 2026

To begin, respondents in the AI interview treatment arms had a substantially higher word count than those who receive the Closed then Fixed open-ended question (CF) (RQ6). Despite the AI interviewer being instructed to ask no more than three follow-up questions,[11] in conditions with the AI interviewer, respondents used 70 [64, 75] more words answering immigration questions and 57 [50, 64] more words answering tariff questions. These rates more than doubled the word counts in CF.

RQ3 asks about the specific difference caused by following up with an AI interviewer, as opposed to beginning the experience with such an encounter. While we saw no such order effect for immigration, we did observe that the Closed, then AI follow-up (CAI) treatment produced significantly fewer words than the other two AI conditions on tariffs. This finding suggests that, at least in some circumstances, the procedure of interviewing prior to asking a closed-ended question encouraged more detailed responses from the respondents.

Moving beyond simple word counts, we analyzed whether there are differences in the reasoning provided by respondents. LLMs are highly proficient in text summarization (Hake et al. 2024; Mathis et al. 2024) and annotation (Savelka and Ashley 2023; Törnberg 2023; Wang et al. 2023). We leveraged these capabilities, in conjunction with human oversight and review, to analyze specific parts of the output of these interviews and the degree to which they were influenced by the use of AI interviewing.[12] We used an LLM, OpenAI's GPT-4o (their flagship model at the time of this study), to summarize the rationales given for respondents' positions on the immigration and tariffs issues, splitting responses into groups that support, oppose, or are neutral on each policy based on closed-ended responses for each treatment.

LLM output is stochastic, yielding different outputs on repeated runs. Therefore, we took a two-step approach to coding. First, for each combination of issue (immigration and tariffs) and position (support, oppose, neutral), we asked the LLM to summarize the "major reasons" provided in the interviews, along with a summary label of their prevalence ("high," "medium," or "low").[13] The results produced an approximately normally distributed count of topics generated, with a median of 7 and a mean of 7.3. Among runs, the maximum number of topics generated across all conditions is 16, and the minimum is 1. Table 3 shows an

11. An instruction which, admittedly, the AI does not always follow; see Supplementary Material section 7.

12. We used the OpenAI API (GPT-4o) for all these analyses. None of the data was stored or used for training by OpenAI; see Supplementary Material section 8.5.

13. Details in Supplementary Material section 9.1.

Downloaded from academic.oup.com/poq/advance-article/doi/10.1093/poq/nfag063/8762750 by guest on 10 November 2026

Table 3. Example of LLM summary of reasons given for support for allowing undocumented workers to stay in the United States. The treatment for this summary is the closed-ended question followed by an AI interview (CAI).

Reason	Summary	Labeled prevalence	Actual prevalence
Economic contribution	Many participants highlight that undocumented immigrants fill essential jobs that US citizens are unwilling to do, thus contributing significantly to the economy.	High	0.437
Family and humanitarian concerns	Participants express concern over family separation and suggest that individuals who have built lives and families in the US should have a path to legality.	High	0.135
Cultural and workforce diversity	Participants acknowledge the role of immigrants in enriching US culture and workforce diversity, adding that immigrants rejuvenate local communities.	Medium	0.067
Long-term residents	Participants argue that undocumented immigrants who have lived in the US for many years and are law-abiding should be allowed to stay legally.	Medium	0.141
Lack of criminal record	Participants frequently cite the absence of a criminal record as a crucial condition for allowing undocumented immigrants to stay legally.	Medium	0.486
Moral and ethical considerations	Some participants talk about moral obligations and empathy for individuals fleeing from dire conditions, suggesting they should be treated humanely.	Medium	0.303

(continued)

Downloaded from academic.oup.com/poq/advance-article/doi/10.1093/poq/nfag063/8762750 by guest on 10 November 2026

Table 3. Continued.

Reason	Summary	Labeled prevalence	Actual prevalence
Legalization process improvements	Participants propose simplifying and accelerating the legal process for undocumented immigrants, as current systems are seen as burdensome and slow.	Low	0.138
Security vetting	A few participants indicate the need for background checks and vetting as reasonable measures to ensure the safety and integrity of the immigration process.	Low	n/a

example of this output for those who support allowing undocumented workers to stay in the United States.[14]

Our second step was to draw a random set of labels from those summaries for each policy and position within the CAI treatment. We then asked the LLM to label comments using these categories.[15] A random sample of 300 labels from each group revealed the labeling to be quite accurate, at least within the tolerance expected from a skilled research assistant.[16]

While the topic summaries generated by the LLM were quite good, as judged by checks by the authors and comparisons against other LLM-based systems,[17] the estimates of reasoning prevalence in table 3 were not. The rightmost two columns report the qualitative estimate of prevalence generated by the LLM and the actual prevalence observed (for the CAI treatment). These columns align for some categories (e.g., "Economic contribution" or "Moral and ethical consideration"), but do not for others (e.g., "Family and humanitarian concerns"). In addition, there appears to be a redundant category. "Security vetting" was not used at all in labeling

14. The illustrated example was simulation 35 from the CAI treatment and was randomly drawn from the simulations for that treatment.

15. Details in Supplementary Material section 9.2.

16. We found only 4 cases (1.3 percent) where the label is clearly incorrect and another 16 (5.3 percent) where it is unclear if the label was correct. Details in Supplementary Material section 9.2.

17. See Supplementary Material section 9.

Downloaded from academic.oup.com/poq/advance-article/doi/10.1093/poq/nfag063/8762750 by guest on 10 November 2026

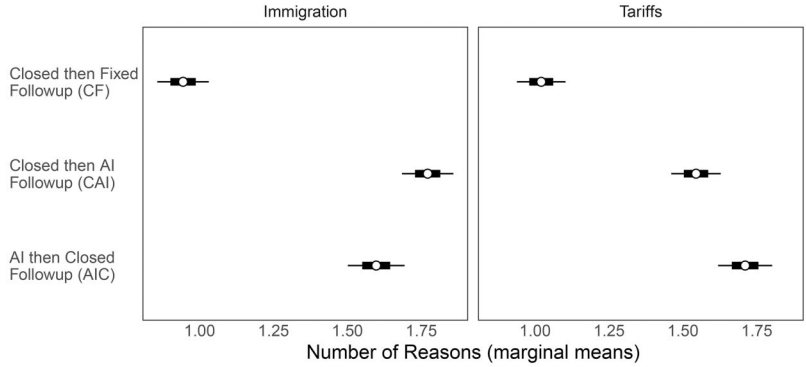


Figure 3. Interaction with the AI interviewer boosted reason provision. Points show marginal means, and bars show 67 percent and 95 percent intervals.

the responses, likely because it was redundant with the "lack of criminal record category." Thus, while the LLM proves effective at summarizing the reasons themselves, it does not do a good job of estimating their prominence or of avoiding nonmutually exclusive categories.

We can now address our research questions RQ3 and RQ6 using this measurement of reason provision. First, we counted the number of reasons each respondent provided separately by issue and regressed the results on treatment indicators (see figure 3). The results were clear: The AI interviewer not only encouraged respondents to write more text, it also caused them to increase the number of reasons they offered. A bivariate regression on an indicator for the AI interviewer revealed that it increases the number of reasons by $0.75\,[95\ \text{percent}\,CIs = 0.65, 0.85]$ for immigration and $0.63\,[0.53, 0.72]$ for tariffs. However, we saw no clear order effect from where the AI was placed. We conclude that AI interviewers sparked longer and more varied conversations than what is achieved by a fixed prompt.

Finally, we estimated separate regression models of binary indicators for whether a respondent provided a reason on treatment condition indicators, excluding the control, Closed Only (CO) condition, which lacked text output, and AI Only (AIO), which lacked a closed-ended interview question, and so could not be binned into support, neutral, or oppose positions. Figure 4 presents the results.[18]

First, we see that in most cases, the Closed, then Fixed follow-up question (CF) yielded a generally lower propensity to provide reasons. Pooling

18. We omitted the neutral position, which has only 161 respondents for tariffs and 66 for immigration.

Downloaded from academic.oup.com/poq/advance-article/doi/10.1093/poq/nfag063/8762750 by guest on 10 November 2026

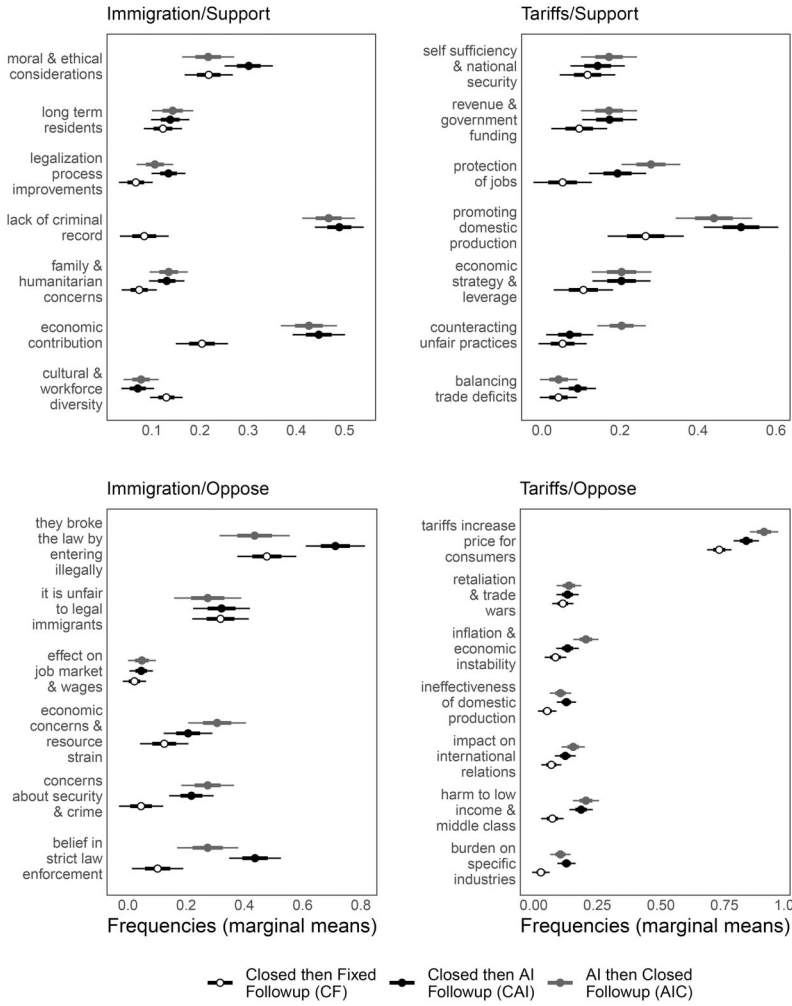


Figure 4. Interaction with the AI interviewer substantively altered reason provision. Points show marginal means, and bars show 67 percent and 95 percent intervals.

across all reasons, both the AIC and CAI conditions had provision rates of 24 percent, while that for CF was only 14 percent. Combined with the above evidence on word counts, the answer to RQ3 appears to be that the AI interviewer both lengthens and broadens engagement.

This difference in engagement also depended on whether the AI interviewer preceded or followed up a closed-ended question (RQ3). At times,

Downloaded from academic.oup.com/poq/advance-article/doi/10.1093/poq/nfag063/8762750 by guest on 10 November 2026

this difference was dramatic. Across all 41 reasons (including the netural positions), the 95 percent interval for the difference between CAI and AIC excluded 0 in 22 percent of cases. Those differences were not univalent—follow-up with the AI interviewer was more likely to cause those who support tariffs to invoke domestic production, while initiating the experience with the AI interviewer was more likely to yield a discussion of job protection. The implication is that the placement of the AI interviewer—whether before or after a closed-ended question—can meaningfully shape the content and depth of the responses, demonstrating how design choices in interview sequencing can influence the substance of the conversation.

The inclusion of AI interviewing had a substantial impact on the topics discussed and the reasoning patterns detected for popular political positions. The interviewing process suggests that these arguments do not emerge by simply asking a follow-up question—a method commonly used in current studies, when open-ended questions are included at all. Further experimental studies are required to uncover the extent to which the arguments preferred are sufficiently important for the extra time and cost of AI interviews.

Effects on Closed-Ended Responses

Beyond participant experience and detail of responses, there are fundamental questions about whether and how AI interviewers might alter closed-ended responses. To shed light on these questions, we embedded the closed-ended items into the Engage interview. We designate these questions as the "Interview" closed-ended items because they were asked as part of the broader interview in Engage. Treatment arms varied based on whether there was no follow-up at all, the Closed Only (CO) control condition, a traditional Closed then Fixed (CF) follow-up, or the closed-ended questions came before (CAI) or after (AIC) interacting with the AI interviewer. We could, therefore, probe for effects of the presence and timing of the AI interviewer by analyzing responses to these closed-ended questions across all four treatment arms. The key contrast was between AIC and the other three treatment arms (RQ2).

Results appear in figure 5. For immigration, interacting with the AI interviewer shifted responses slightly toward greater support. A bivariate regression of responses to this question on an indicator for treatment arm AIC yielded a coefficient of $0.35\,[0.13, 0.57]$. In contrast, we saw no such effect for tariffs. These findings suggest that interacting with the AI interviewer can move responses, but that such effects are likely to be context specific. For some issues, discussion with an AI interviewer prior to answering a closed-ended survey question may influence responses,

Downloaded from academic.oup.com/poq/advance-article/doi/10.1093/poq/nfag063/8762750 by guest on 10 November 2026

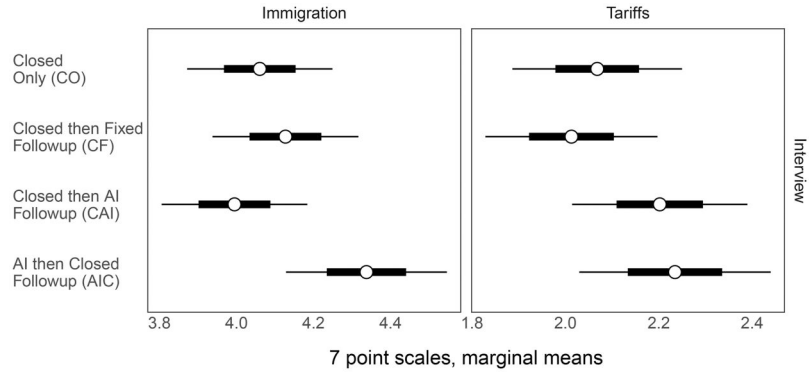


Figure 5. Interacting with the AI interviewer sometimes alters responses to closed-ended questions. Points show marginal means, and bars show 67 percent and 95 percent intervals.

but, for other issues, it seems to have no such effect. The factors underlying this may form an interesting area for future study.

After respondents finished the interview on the Engage platform, the survey continued on Qualtrics. After a series of personality and political questions, we fielded the same closed-ended items again. We used responses to these "Post-Interview" questions to gauge the extent to which the AI interviewer provoked change (RQ4). To do this, we first modeled a binary indicator capturing any change between Interview and Post-Interview responses, and second, we analyzed the magnitude of change by measuring the absolute difference between the two.

Figure 6 displays the results. Across both issues, we saw evidence that interacting with the AI interviewer *after* providing a response to the closed-ended question (CAI) caused some respondents to change. There is inconsistent evidence that *any* follow-up probe for explanation caused such changes—the difference between "any change" rates for the Closed, then Fixed (CF) probe and control Closed Only (CO) condition was $0.06\,[0.02, 0.11]$ for immigration and $0.03\,[-0.01, 0.07]$ for tariffs. But the corresponding contrasts between CAI and CO were $0.13\,[0.09, 0.18]$ and $0.08\,[0.04, 0.11]$. Moreover, even though collapsing all changes into this binary measure entailed eliminating any information about the size of these changes, this pattern persisted in the analysis of the distance between Interview and Post-Interview responses.

Given that the AI interviewer provoked changes, we next analyze the extent to which participants selected the most extreme response options (RQ5). To do so, we calculated the distance of responses to the scale median, then subtracted the distance observed during the Interview from that

Downloaded from academic.oup.com/poq/advance-article/doi/10.1093/poq/nfag063/8762750 by guest on 10 November 2026

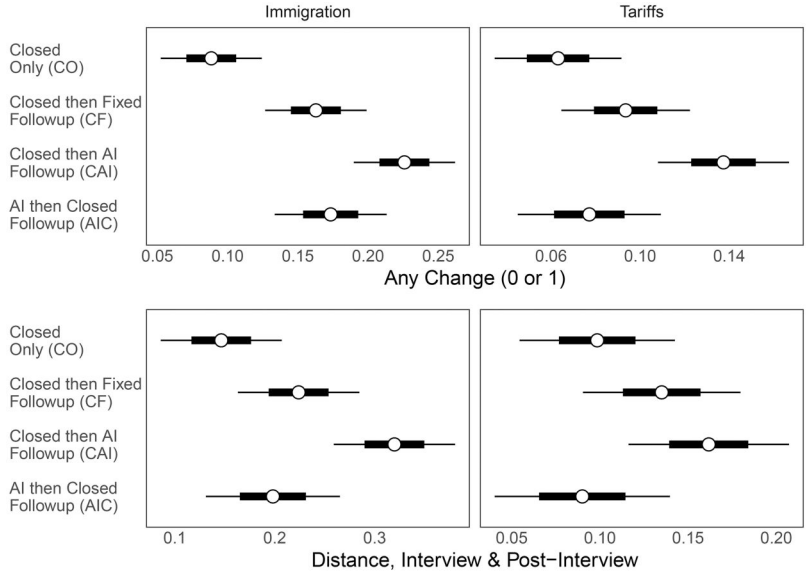


Figure 6. Follow-up with the AI interviewer changed responses. Points show marginal means, and bars show 67 percent and 95 percent intervals.

during Post-Interview, and then regressed that quantity on the treatment arm indicators.

Results are depicted in figure 7. We observed that—relative to the control Closed Only (CO) condition—closed-ended questions that are followed up by the AI interviewer (CAI) polarized responses on both issues. For immigration, the effect was $0.064\,[0.007, 0.123]$, and for tariffs, $0.045\,[0.02, 0.090]$. This effect was not limited to the AI follow-up, however, as we saw similar changes for follow-up with a fixed open-ended prompt (CF, $0.057\,[-0.003, 0.114]$ for immigration and $0.024\,[-0.018, 0.066]$ for tariffs). These effects were of small magnitude. The outcome variable has a theoretical range from -3 to 3, and, in keeping with the modest change rates visible in figure 6, most of its values are zero. Nevertheless, these analyses reveal that, when follow-up provoked change, it tended toward response polarization.

This evidence suggests that AI interviewers incite several changes in responses. First, conducting a preliminary AI interview impacts responses on immediate follow-up, closed-ended questions in some cases. Second, reversing that order, following up a closed-ended item with a probing AI interview makes changes in subsequent answers to the same question more likely, even beyond changes caused by fixed follow-up prompts. Finally,

Downloaded from academic.oup.com/poq/advance-article/doi/10.1093/poq/nfag063/8762750 by guest on 10 November 2026

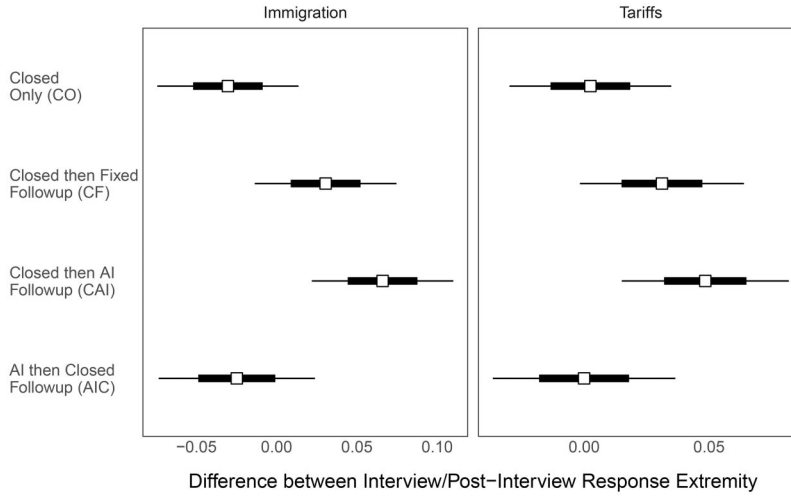


Figure 7. Follow-up probes polarized responses. We calculate the outcome variable by taking the distance of responses to the scale median, and subtracting the distance observed during the interview from that during Post-Interview. Points show marginal means, and bars show 67 percent and 95 percent intervals.

while these changes are relatively uncommon, they tend to push toward extreme response options. These results suggest that some level of "deliberation within" or response opinion justification may be triggered by the use of AI interviewing, altering the cognitive process respondents use when formulating responses to closed-ended questions and, perhaps, crystallizing attitudes.

Conclusions

We present a large-scale randomized survey experiment varying both the presence and timing of an AI interviewer to investigate its effects across multiple dimensions. A central finding is that AI interviewing does not negatively affect item nonresponse (rolloff) or participant satisfaction, when compared against exclusively closed-ended or fixed open-ended follow-up questions, despite increasing interview duration and cognitive engagement. AI interviewing substantially enhances both the quantity and depth of respondents' articulated reasoning compared with non-AI counterparts, suggesting improved access to the underlying, internally reflective deliberation informing public opinion. However, the AI interviewer elicits different responses than traditional methods.

Downloaded from academic.oup.com/poq/advance-article/doi/10.1093/poq/nfag063/8762750 by guest on 10 November 2026

While fostering more explicit reasoning, this increased engagement may induce cognitive shifts, evidenced by changes in subsequent closed-ended responses, including signs of opinion polarization. In this sense, similar to findings with "think aloud" interview procedures, AI may act as an intervening stimulus—amplifying the salience of particular considerations and prompting a more deliberative reasoning process.

In terms of the AI interview implementation, we found some difference in duration of surveys, with respondents taking more time when a closed-ended question followed an AI interview than in the case of either the AI interview following the closed-ended question or the AI interview only. This difference in time spent, however, did not produce a consistent difference across topics in terms of the number of reasons provided for their answers, nor did it produce substantial differences in the types of reasons offered. We did note, however, that following a closed-ended question with an AI interview was more likely than other conditions to change and polarize subsequent responses to the same question.

Further research should explore the cognitive mechanisms through which AI-mediated interviews influence the salience and weighting of considerations during opinion formation and assess the extent to which these effects generalize across diverse issues and populations. For example, questions on less polarizing issues, issues not related to politics, or involving sensitive topics may alter respondent behavior in AI interviews. Additionally, researchers should examine how AI interviewer characteristics—such as affect, tone, and conversational structure—impact respondent behavior, reasoning, and overall interview quality.[19] Further, researchers should consider constructs, such as perceived warmth, competence, comfort, and trust in the AI interviewers, from respondents to evaluate under what conditions respondents are more (or less) willing and able to engage with these tools. Understanding these dynamics is critical for evaluating the long-term implications of AI interviewing on the reliability, stability, and comparability of public opinion measures across time and contexts.

Appendix A. AAPOR-Required Disclosure Elements

Data Collection Strategy

Survey.

19. For more detailed analysis of the AI interviewer behavior in this study, see Supplementary Material section 7.

Downloaded from academic.oup.com/poq/advance-article/doi/10.1093/poq/nfag063/8762750 by guest on 10 November 2026

Who Sponsored the Research and Who Conducted It

The survey was conducted through CloudResearch's Connect platform. The authors undertook all analyses. There was no sponsor besides the authors.

Measurement Tools/Instruments

Questionnaires with survey questions, response options, and the AI prompts are reported in Supplementary Material section 2.

Population Under Study

US citizens age 18 or older.

Method Used to Generate and Recruit the Sample

CloudResearch recruits respondents through convenience sampling. Respondents sign up to take surveys online through CloudResearch's interface. We chose to have CloudResearch match the age, gender, ethnicity, and race characteristics of our sample against US Census values. Respondents were paid $3 for their participation, delivered through CloudResearch.

Methods and Modes of Data Collection

Surveys were administered online on respondents' computers or tablets (phones were not allowed) on the Qualtrics platform. AI interviews were conducted using the Engage survey system.

Dates of Data Collection

December 6–11, 2024.

Sample Size

2,243 respondents.

How the Data Were Processed and Procedures to Ensure Data Quality

We embedded an attention check (see Supplementary Material section 2) within the survey. Those who failed the attention check were not removed from the survey pool, but results were estimated both with and

Downloaded from academic.oup.com/poq/advance-article/doi/10.1093/poq/nfag063/8762750 by guest on 10 November 2026

without those who failed the attention check (see Supplementary Material section 10).

Dispositions Rates

2,250 completed, 33 timed out, 73 returned.

Final Sample Size

2,243, due to missing demographic data from CloudResearch.

Measurement and Model Specification

All results are estimated using OLS in R. Details of AI text coding procedures and checks can be found in Supplementary Material section 9.

General Statement Acknowledging Limitations of the Design and Data Collection

Because it is an online convenience sample, this sample may differ from the US population on dimensions correlated with panel membership or survey-taking propensity. In addition, we are limited to testing one AI interview system. Other systems may perform differently and evoke different reactions from respondents.

Supplementary Material

Supplementary Material may be found in the online version of this article: https://doi.org/10.1093/poq/nfag063.

Acknowledgments

The authors would like to thank the discussants and audience members at the 2025 annual meeting of the Midwest Political Science Association (MPSA) in Chicago, IL, and the American Association for Public Opinion Research (AAPOR) in St. Louis, MO. Authors are presented in alphabetical order.

Data Availability

The data underlying this article are available at https://doi.org/10.7910/DVN/DA2URD.

Downloaded from academic.oup.com/poq/advance-article/doi/10.1093/poq/nfag063/8762750 by guest on 10 November 2026

References

Barari, S., J. Angbazo, N. Wang, et al. 2025. "AI-Assisted Conversational Interviewing: Effects on Data Quality and Respondent Experience." arXiv preprint arXiv : 2504.13908, preprint: not peer reviewed.

Brooks, A. W., and L. K. John. 2018. "The Surprising Power of Questions." *Harvard Business Review* 96:60–67.

Cacioppo, J. T., and R. E. Petty. 1982. "The Need for Cognition." *Journal of Personality and Social Psychology* 42:116–31.

Chopra, F., and I. Haaland. "Conducting Qualitative Interviews with AI." 2023. SSRN Pre-print. https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4583756

Croes, E. A., and M. L. Antheunis. 2020. "36 Questions to Loving a Chatbot: are People Willing to Self-Disclose to a Chatbot?" In *International Workshop on Chatbot Research and Design*, edited by A. Folstad, 81–95. Springer.

Daikeler, J., M. Bošnjak, and K. Lozar Manfreda. 2020. "Web versus Other Survey Modes: An Updated and Extended Meta-Analysis Comparing Response Rates." *Journal of Survey Statistics and Methodology* 8:513–39.

Douglas, B. D., P. J. Ewell, and M. Brauer. 2023. "Data Quality in Online Human-Subjects Research: Comparisons between Mturk, Prolific, Cloudresearch, Qualtrics, and Sona." *Plos One* 18:e0279720.

Feldman, S. 1989. "Measuring Issue Preferences: The Problem of Response Instability." *Political Analysis* 1:25–60.

Galesic, M. 2006. "Dropouts on the Web: Effects of Interest and Burden Experienced during an Online Survey." *Journal of Official Statistics* 22:313.

Geiecke, F., and X. Jaravel. 2024. "Conversations at Scale: Robust AL-Led Interviews with a Simple Open-Source Platform." SSRN Working Paper No. 4974382.

Goodin, R. E. 2000. "Democratic Deliberation within." *Philosophy & Public Affairs* 29:81–109.

Goodin, R. E., and S. J. Niemeyer. 2003. "When Does Deliberation Begin? internal Reflection versus Public Discussion in Deliberative Democracy." *Political Studies* 51:627–49.

Hake, J., M. Crowley, A. Coy, et al. 2024. "Quality, Accuracy, and Bias in Chatgpt-Based Summarization of Medical Abstracts." *Annals of Family Medicine* 22:113–20.

Hardin, C. D., and E. T. Higgins. 1996. "Shared Reality: How Social Verification makes the Subjective Objective." In *Handbook of Motivation and Cognition*, edited by R. M. Sorrentino and E. T. Higgins, Vol. 3, 28–84. The Guilford Press.

Hart, E., E. M. VanEpps, and M. E. Schweitzer. 2021. "The (Better than Expected) Consequences of Asking Sensitive Questions." *Organizational Behavior and Human Decision Processes* 162:136–54.

Hartman, E., and F. D. Hidalgo. 2018. "An Equivalence Approach to Balance and Placebo Tests." *American Journal of Political Science* 62:1000–1013.

Hartman, R., A. J. Moss, S. N. Jaffe, C. Rosenzweig, L. Litman, and J. Robinson. 2023. "Introducing Connect by Cloudresearch: Advancing Online Participant Recruitment in the Digital Age." *PsyArXiv*, preprint: not peer reviewed. https://osf.io/preprints/psyarxiv/ksgyr_v1?utm_source=hs_email&utm_medium=email&_hsenc=p2ANqtz-8FJhys54lcxx3QausORM8iReCMVXQVa-lTNRsSoZWNBp-lSY9jKdrlVd3wI2QglhY_ZFqQ

Hauser, D. J., A. J. Moss, C. Rosenzweig, S. N. Jaffe, J. Robinson, and L. Litman. 2023. "Evaluating Cloudresearch's Approved Group as a Solution for Problematic Data Quality on Mturk." *Behavior Research Methods* 55:3953–64.

Heerwegh, D., and G. Loosveldt. 2008. "Face-to-Face versus Web Surveying in a High-Internet-Coverage Population: Differences in Response Quality." *Public Opinion Quarterly* 72:836–46.

Downloaded from academic.oup.com/poq/advance-article/doi/10.1093/poq/nfag063/8762750 by guest on 10 November 2026

Herzberg, P. Y., and E. Brähler. 2006. "Assessing the Big-Five Personality Domains via Short Forms." *European Journal of Psychological Assessment* 22:139–48.

Kalla, J. L., and D. E. Broockman. 2020. "Reducing Exclusionary Attitudes through Interpersonal Conversation: Evidence from Three Field Experiments." *American Political Science Review* 114:410–25.

King, G., M. Tomz, and J. Wittenberg. 2000. "Making the Most of Statistical Analyses: Improving Interpretation and Presentation." *American Journal of Political Science* 44:347–61.

Kreuter, F., S. Presser, and R. Tourangeau. 2008. "Social Desirability Bias in Cati, Ivr, and Web Surveys: The Effects of Mode and Question Sensitivity." *Public Opinion Quarterly* 72:847–65.

Lazarsfeld, P. F. 1944. "The Controversy over Detailed Interviews—an Offer for Negotiation." *Public Opinion Quarterly* 8:38–60.

Lodge, M., and C. S. Taber. 2013. *The Rationalizing Voter*. Cambridge University Press.

Manfreda, K. L., M. Bosnjak, J. Berzelak, I. Haas, and V. Vehovar. 2008. "Web Surveys versus Other Survey Modes: A Meta-Analysis Comparing Response Rates." *International Journal of Market Research* 50:79–104.

Mathis, W. S., S. Zhao, N. Pratt, J. Weleff, and S. De Paoli. 2024. "Inductive Thematic Analysis of Healthcare Qualitative Interviews Using Open-Source Large Language Models: How Does It Compare to Traditional Methods?" *Computer Methods and Programs in Biomedicine* 255:108356.

Minozzi, W., R. Kennedy, K. M. Esterling, M. A. Neblo, and R. Jewell. 2024. "Testing the Benefits of Public Deliberation." *American Journal of Political Science* 68:319–34.

Pew Research Center. 2018. "As New Tariffs Take Hold, More See Negative than Positive Impact for the U.S." Date accessed: 2025-05-01.

———. 2024. "Most Americans Say Undocumented Immigrants Should be Able to Stay Legally Under Certain Conditions." Date accessed: 2025-05-01.

Reja, U., K. L. Manfreda, V. Hlebec, and V. Vehovar. 2003. "Open-Ended vs. close-Ended Questions in Web Questionnaires." *Developments in Applied Statistics* 19:159–77.

Savelka, J., and K. D. Ashley. 2023. "The Unreasonable Effectiveness of Large Language Models in Zero-Shot Semantic Annotation of Legal Texts." *Frontiers in Artificial Intelligence* 6:1279794.

Siena College Research Institute. 2024. "Who is Kamala Harris?" Date accessed: 2025-06-11.

Taber, C. S., and M. Lodge. 2006. "Motivated Skepticism in the Evaluation of Political Beliefs." *American Journal of Political Science* 50:755–69.

Thompson, D. F. 2008. "Deliberative Democratic Theory and Empirical Political Science." *Annual Review of Political Science* 11:497–520.

Törnberg, P. 2023. "Chatgpt-4 Outperforms Experts and Crowd Workers in Annotating Political Twitter Messages with Zero-Shot Learning." arXiv Preprint arXiv : 2304.06588, preprint: not peer reviewed.

Tourangeau, R., L. J. Rips, and K. Rasinski. 2000. *The Psychology of Survey Response*. Cambridge University Press.

Tourangeau, R., and T. Yan. 2007. "Sensitive Questions in Surveys." *Psychological Bulletin* 133:859–83.

Tsai, L. L., A. Pentland, A. Braley, N. Chen, J. R. Enríquez, and A. Reuel. 2024. "Generative AI for Pro-Democracy Platforms." *An MIT Exploration of Generative AI*. https://mit-genai.pubpub.org/pub/mn45hexw/release/1

Vedejová, D., and V. Čavojová. 2022. "Confirmation Bias in Information Search, Interpretation, and Memory Recall: Evidence from Reasoning about Four Controversial Topics." *Thinking & Reasoning* 28:1–28.

Velez, Y. R., and P. Liu. 2025. "Confronting Core Issues: A Critical Assessment of Attitude Polarization Using Tailored Experiments." *American Political Science Review* 119:1036–53.

Downloaded from academic.oup.com/poq/advance-article/doi/10.1093/poq/nfag063/8762750 by guest on 10 November 2026

Wang, Z., Y. Pang, and Y. Lin. 2023. "Large Language Models are Zero-Shot Text Classifiers." arXiv preprint arXiv : 2312.01044, preprint: not peer reviewed.

Willis, G. B. 2004. *Cognitive Interviewing: A Tool for Improving Questionnaire Design*. Sage Publications.

Wuttke, A., M. Aßenmacher, C. Klamm, M. M. Lang, Q. Würschinger, and F. Kreuter. 2024. "AI Conversational Interviewing: Transforming Surveys with LLMS as Adaptive Interviewers." arXiv preprint arXiv : 2410.01824, preprint: not peer reviewed.

Xiao, Z., M. X. Zhou, and Q. V. Liao, et al. 2020. "Tell Me about Yourself: Using an AI-Powered Chatbot to Conduct Conversational Surveys with Open-Ended Questions." *ACM Transactions on Computer-Human Interaction (TOCHI)* 27:1–37.

Yin, Y., N. Jia, and C. J. Wakslak. 2024. "AI Can Help People Feel Heard, but an AI Label Diminishes This Impact." *Proceedings of the National Academy of Sciences of the United States of America* 121:e2319112121.

Zaller, J. 1992. *The Nature and Origins of Mass Opinion*. Cambridge University Press.

© The Author(s) 2026. Published by Oxford University Press on behalf of American Association for Public Opinion Research.
This is an Open Access article distributed under the terms of the Creative Commons Attribution License (https://creativecommons.org/licenses/by/4.0/), which permits unrestricted reuse, distribution, and reproduction in any medium, provided the original work is properly cited.
Public Opinion Quarterly, 2026, 00, 1–29
https://doi.org/10.1093/poq/nfag063
Original article

Downloaded from academic.oup.com/poq/advance-article/doi/10.1093/poq/nfag063/8762750 by guest on 10 November 2026